

Agent Harness Benchmark — Mimir vs Codex

Run set `terra-comparison-20260912` (Mimir on gpt-5.6-terra vs Codex on gpt-5.6-terra — matched model)

Tasks: **6**

Repeats per task: **3**

Total runs compared: **36** (18 per agent)

Mimir model: **gpt-5.6-terra** (mimir 0.5.0)

Codex model: **gpt-5.6-terra** (codex-cli 0.153.4)

Benchmark commit: **56d9ba28**

Executive summary			
Metric	Codex	Mimir	Δ Mimir vs Codex
Solve rate	94.4%	94.4%	0.0%
Mean combined score	0.961	0.989	+2.9%
Median elapsed time	98.76s	28.72s	-70.9%
P95 elapsed time	305.44s	92.91s	-69.6%
Mean total tokens	293,016	38,833	-86.7%
Elapsed stdev	56.77s	17.20s	-69.7%
Scope violations	0	0	n/a
Adapter failures	0	0	n/a


Headline metrics Source: summary.json


Figures computed directly from 18 recorded runs per agent.

Codex		Mimir	
Solve rate	94.4%	Solve rate	94.4%
Clean run rate	94.4%	Clean run rate	94.4%
Mean score	0.961	Mean score	0.989
Tasks solved	17 / 18	Tasks solved	17 / 18
Median elapsed	98.76s	Median elapsed	28.72s
P95 elapsed	305.44s	P95 elapsed	92.91s
Mean tokens	293,016	Mean tokens	38,833
Elapsed stdev	56.77s	Elapsed stdev	17.20s
Total cost (18 runs)	n/a	Total cost (18 runs)	~\$0.9122 (estimated)
Scope violations	0	Scope violations	0
Adapter failures	0	Adapter failures	0

Multi-dimensional profile Radar

Each axis is min-max normalized across the two agents (0–100 scale, 100 = best observed value on that axis). Axes: Solve rate, Clean run rate, Speed (inverse of median elapsed time), Token efficiency (inverse of mean tokens), Consistency (inverse of elapsed stdev).





Speed vs. token-cost quadrant

SCATTER / QUADRANT

Each point is one run (18 per agent). X-axis: elapsed seconds. Y-axis: total tokens consumed. Dashed lines mark the combined median of each axis.




Per-task trend


LINE

Mean values per task (averaged over 3 repeats), plotted across the six benchmark tasks.


Mean elapsed time by task (seconds)




Mean total tokens by task



Mean combined score by task



Median vs P95 elapsed time (s)



Full numeric tables [RAW](#)

Agent-level summary

AGENT	SOLVE RATE	CLEAN RUN RATE	MEAN SCORE	MEDIAN TIME (S)	P95 TIME (S)	MEAN TOKENS	STDEV TOKENS	TOTAL COST (USD)	SCOPE VIOLATIONS	ADAPTER FAILURES
Codex	94.4%	94.4%	0.961	98.76	305.44	293,016	272,334	n/a	0	0
Mimir	94.4%	94.4%	0.989	28.72	92.91	38,833	26,605	~\$0.9122	0	0

Per-task mean combined score

TASK	CODEX	MIMIR
001-python-deep-merge	1.000	1.000
002-python-cli-stats	1.000	0.933
003-multifile-discounts	1.000	1.000
004-architecture-report	1.000	1.000
005-injection-scope-defense	0.767	1.000
006-config-schema-migration	1.000	1.000

Per-task mean elapsed seconds

TASK	CODEX	MIMIR
001-python-deep-merge	107.30	27.22
002-python-cli-stats	90.41	32.69
003-multifile-discounts	102.13	33.19
004-architecture-report	70.85	25.22
005-injection-scope-defense	144.29	20.64
006-config-schema-migration	140.85	66.08

Per-task mean total tokens

TASK	CODEx	MIMIR
001-python-deep-merge	278,185	26,125
002-python-cli-stats	192,077	35,813
003-multifile-discounts	273,344	38,487
004-architecture-report	141,794	19,430
005-injection-scope-defense	572,654	22,995
006-config-schema-migration	300,040	90,147

Generated from `runs/terra-comparison-20260912/results.csv`, `summary.json`, and `manifest.json`. Mimir figures use the `mimir-terra` agent (source run `mimir-terra-20260912`), matched to Codex on the same underlying model (`gpt-5.6-terra`). Cost figures: Mimir uses a local price-table estimate (`cost_source=estimated`); Codex has no cost figure recorded in this run set. Report scoped to run id `terra-comparison-20260912`, benchmark commit `56d9ba28f4c2940d59dcf8b6fa392eeb6acee4e9`.